

CS541 Artificial Intelligence


Guest Lecture on Mean Estimation

Lecturer: Shiwei Zeng


szeng4@stevens.edu

Estimating Average Height

- Assume $D = N(60, 1)$



- Assume $E[D] = 60, \text{Var}[D] = 1$



- Estimator $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$ $S \sim \mathcal{D}^m$


$$E[\hat{\mu}] = E_{S \sim \mathcal{D}^m} \left[\frac{1}{n} \sum_{i=1}^n x_i \right]$$

$$= \frac{1}{n} \sum_{i=1}^n \underbrace{E[x_i]} = 60 = \text{ground truth.}$$

$$\text{Var}_{S \sim \mathcal{D}^m}[\hat{\mu}] = \frac{1}{n^2} \sum_{i=1}^n \text{Var}[x_i] = \frac{1}{n^2} \cdot n = \frac{1}{n} \quad n \uparrow$$



ME in Higher Dimension



D

$$E[D] = ?$$

When Data is Noisy

Adversary: corrupt. ϵ -fraction, $\epsilon < \frac{1}{2}$

Total variation distance D_1, D_2
 $\frac{1}{2} \int |\phi_1 - \phi_2| dx = \frac{\epsilon}{1-\epsilon}$

- 1-dimensional: (a lower bound)

$\phi_1 \leftarrow D_1 = N(\mu_1, 1)$

$D_2 = N(\mu_2, 1) \nearrow \phi_2$

$|\mu_1 - \mu_2| \geq \Omega(\epsilon)$

Q_1, Q_2

$D_\epsilon = (1-\epsilon)D_1 + \epsilon \cdot Q_1 = (1-\epsilon)D_2 + \epsilon Q_2$

$Q_1 = \frac{1-\epsilon}{\epsilon} (\phi_2 - \phi_1) \cdot \mathbb{1}_{\phi_2 \geq \phi_1}$

$Q_2 = \frac{1-\epsilon}{\epsilon} (\phi_1 - \phi_2) \cdot \mathbb{1}_{\phi_1 \geq \phi_2}$


Verify:

$q_1 = (1-\epsilon) \cdot \phi_1 + \frac{1-\epsilon}{\epsilon} (\phi_2 - \phi_1) \cdot \mathbb{1}_{\phi_2 \geq \phi_1}$
 $= \begin{cases} (1-\epsilon) \cdot \phi_2 & \phi_2 \geq \phi_1 \\ (1-\epsilon) \cdot \phi_1 & \phi_2 < \phi_1 \end{cases}$

$q_2 = (1-\epsilon)\phi_2 + \epsilon \frac{1-\epsilon}{\epsilon} (\phi_1 - \phi_2)$
 $= \begin{cases} \end{cases}$

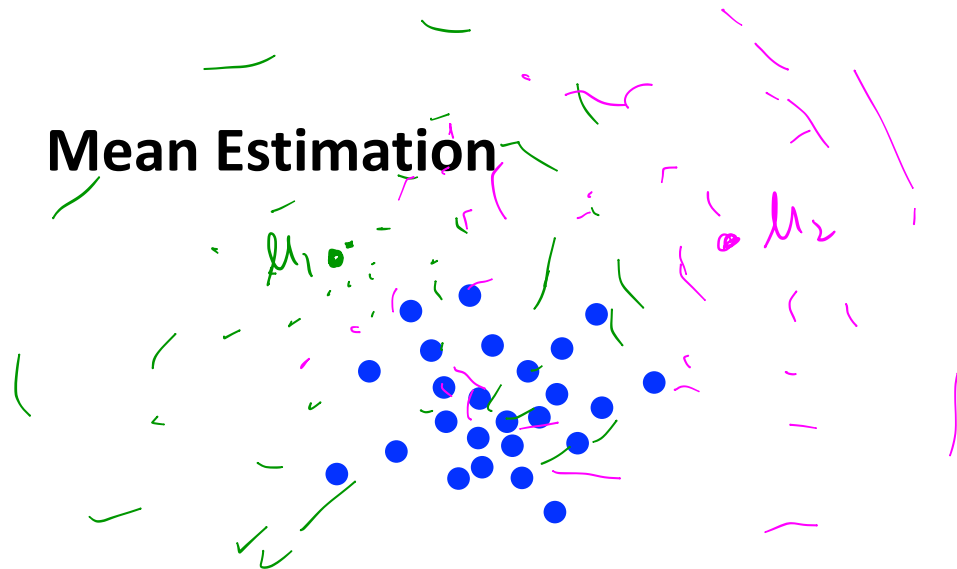
$$D_E = (1-\epsilon) D_1 + \epsilon Q_1$$

$$D_E = (1-\epsilon) D_2 + \epsilon Q_2$$



Robust Mean Estimation

Mean Estimation




D

$$E[D] = ?$$

$$[\frac{1}{2}, 1)$$

ϵ -robust Mean Estimation



An ϵ fraction is corrupted

$D + D'$

$$E[D] = ?$$

$$\Sigma(\xi)$$

bf 2016

$$\underbrace{O(\epsilon \cdot \sqrt{n})}_{1000 \cdot \epsilon}$$

$$n = 10^6$$

$$\Sigma = (0, \frac{1}{2})$$

$$\| \hat{\mu} - \mu \|_2 \geq 500$$

Natural approaches

- Learn each coordinate separately

$$|\hat{\mu} - \mu| \geq \Omega(\varepsilon)$$

in n -dimension : $\|\hat{\mu} - \mu\|_2^2 = \sum_{i=1}^n \underbrace{|\hat{\mu}_i - \mu_i|^2}_{\geq \Omega(\varepsilon)^2} \geq n \cdot \Omega(\varepsilon)^2 \geq \Omega(n\varepsilon^2)$

$$\underline{\Omega(\sqrt{n} \cdot \varepsilon)}$$

$$e^n$$

Natural approaches

- Maximum Likelihood Estimator

Negative Log likelihood = NLL

$$\min \text{NLL}(F, x_1, \dots, x_m) = - \sum_{i=1}^m \log F(x_i)$$

$$F(x_i) = \frac{1}{\sqrt{2\pi}} e^{-\frac{\|x_i - \mu\|_2^2}{2}} \quad \leftarrow \text{Var} = 1$$

$$\min - \sum_{i=1}^m \log \left(\frac{1}{\sqrt{2\pi}} \cdot e^{-\frac{\|x_i - \mu\|_2^2}{2}} \right)$$

$$\Rightarrow \min_{\mu \in \mathbb{R}^n} - \sum_{i=1}^m \left(\log \left(\frac{1}{\sqrt{2\pi}} \right) - \frac{\|x_i - \mu\|_2^2}{2} \right)$$

$$\Rightarrow \arg \min \frac{1}{2} \sum_{i=1}^m \|x_i - \mu\|_2^2 = \hat{\mu} \quad \leftarrow \text{empirical mean.}$$

↓

Can be quite bad.

Alg is not robust.

~ 7:30

Efficient Algorithm – Convex Programming

weight vector $\hat{\mathbf{w}} = (w_1, w_2, \dots, w_m)$

Goal: output $\hat{\mathbf{w}}$, $\sum_{i=1}^m \hat{w}_i \cdot \chi_i = \hat{\mu} \rightarrow \mu$.

min empirical variance.

s.t. $\hat{\mathbf{w}} \in \mathcal{W} \leftarrow O(n^6)$

$$p(x) = x \cdot v^*$$



$$O(\sqrt{n})$$

$$D = \mathcal{N}(\mu, I_n)$$

Efficient Robust Mean Estimation - Filter

1. Compute empirical mean and covariance μ_T, Σ_T *T : corrupted data set.*
2. Compute largest eigenvalue $\underline{\lambda^*}$ of $\Sigma_T - I$, and eigenvector $\underline{v^*}$
3. If $\underline{\lambda^*}$ is small, return $\underline{\mu_T}$ $\underline{\lambda^*} \Sigma = \underline{v^*} -$
4. Otherwise, find $t > \underline{C_1}$ such that \leq

$$\Pr_{X \in T} [\underline{|\underline{v^*} \cdot (X - \underline{\mu_T})|} > t] > C_2 e^{-t^2/2} + \frac{C_3 \varepsilon}{t^2 \log(n \log \frac{n}{\varepsilon T})}$$
5. Remove X such that $|\underline{v^*} \cdot (X - \underline{\mu_T})| > \underline{t}$, go back to step 1.

$$O(\varepsilon)$$

λ^* : eigenvalue \longrightarrow variance.


$$\lambda^* v^* = v^* \Sigma_T$$

$$\begin{aligned} \boxed{\text{Var}[x \cdot v^*]} &= E[(x \cdot v^*)^2] \\ x &\sim N(0, I) \\ &= E_{x \sim \mathcal{D}}[(v^{*T} x)(x^T v^*)] \\ &= v^* \cdot \underbrace{E[x x^T]}_{= \Sigma_T} \cdot v^* \\ &= v^* \cdot \Sigma_T \cdot v^* \\ &= \lambda^* v^* \cdot v^* \xrightarrow{=1} \\ &= \lambda^* \end{aligned}$$

$$\epsilon \geq \frac{1}{2} \quad \alpha = 1 - \epsilon$$

List-decodable Mean Estimation

Mean Estimation




D

$$E[D] = ?$$

List-Decodable Mean Estimation

$$10^6 \times 20\% = 0.2 \times 10^6 = 200000$$



$$E[D] \in \{\mu_1, \dots, \mu_m\}$$

Gaussian Annulus Theorem.

Algorithm: Multi-filtering

- A tree of subsets T_i 's, ① Clustering

- ② Iterate through each node


(1) Create a leaf node, an estimate $\hat{\mu}_i$

(2) Create child nodes, subsets T_i 's

- One node, cleaner set
- Two nodes, overlapping subsets

(3) Delete if it can't be α -good.

- No more filtering, then return all $\hat{\mu}_i$'s




$$T_1 \cap T_2 = \emptyset$$

α -fraction

$$E[XX^T]$$

$$\begin{bmatrix} x_1 & x_2 & x_3 & \dots \\ x_1 & x_2 & x_3 & \dots \\ \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$



$$O\left(\log \frac{1}{\alpha}\right)$$

$$O\left(\frac{1}{\alpha^{\frac{1}{2d}}}\right) \quad d \in \mathbb{N}^+, \quad d: \text{degree of polynomial}$$

$$T_i: \alpha T_i \rightarrow \text{good samples.}$$

$$\text{err} = \min_i \|\hat{\mu}_i - \mu\|_2 \leq O\left(\frac{1}{\sqrt{\alpha}}\right)$$